

Harvard University: Data Science Capstone Report

Ferry Edouard

2023-11-27

Summary

The purpose of this project is to predict movie ratings using the MovieLens(10M) dataset, which contains the ratings of many movies given by many users. For this project we begin by importing (downloading) data, cleaning and preparation of the EDA(exploratory data analysis). Hence we explored the dataset in order to get valuable outcomes. The dataset is divided into a 90/10 (ratio) split of training and test datasets. A model is built from the training data, and then applied to the unseen test data. The success metric is root mean square estimate (RMSE). We attempted to predict movie ratings on a scale of 0.5 stars to 5 stars, and RMSE is measured on the same scale. An RMSE of 0 means we are always correct, which is unlikely. An RMSE of 1 means the predicted ratings are off by 1 star. Note that the `final_holdout_test` set is used only for the final RMSE evaluation. The goal for this project is to achieve $RMSE < 0.8649$ as computed on the unseen test dataset.

Introduction

Nowadays recommendation systems are omnipresent and trained to understand the preferences, previous decisions, and characteristics of people and products using data gathered about their interactions. Recommender systems are highly useful as they help users discover products and services they might otherwise have not found on their own. Recommendation systems can also be as: identifying and pursuing key business channels(selling products based upon customer's specific needs/expectations and so on and so forth). Hence (in this vein) comes to mind (life) the idea of building a movie recommendation system: the MovieLens data set as a project using the 10M version, a data set that was collected by GroupLens Research.

Project Goal:

The purpose of choosing MovieLens data set is to predict movie ratings based on user preference, age of movie, genre/category of movies etc... This will be done by training a machine learning algorithm that predicts user ratings (on a scale of 0.5 to 5 stars) using the MovieLens dataset split into training and validation sets to train on and predict movie ratings the validation set.

Note that the `final_holdout_test` set is used only for the final RMSE evaluation. The goal for this project is to achieve $RMSE < 0.8649$ as computed on the unseen test dataset.

The first step toward this goal will be to look at the schema/structure of the data set, visualize it and then sequentially (progressively) build a model that will meet the expectations (reach target accurately). Hence, begins my journey.

This project is the last for *Data Science: Capstone* (PH125.9x) course in the HarvardX Professional Certificate Data Science Program. We will be using the methods taught in the program.

Data exploration

```
names(edx)
```

```
## [1] "userId" "movieId" "rating" "timestamp" "title" "genres"
```

```
### The following questions are part of the Quiz which is ###  
### 10% of the project. I'll make sure no duplicate answers #####
```

```
# How many zeroes & 3s
```

```
zeros <- sum(edx$rating == 0)  
threes <- sum(edx$rating == 3)
```

```
## Q2
```

```
# How many zeros were given as ratings in the edx dataset?  
# How many threes were given as ratings in the edx dataset?
```

```
print(zeros)
```

```
## [1] 0
```

```
print(threes)
```

```
## [1] 2121240
```

```
# Q3  
# How many different movies are in the edx dataset?  
uniqueMovies <- length(unique(edx$movieId))  
n_distinct(edx$movieId)
```

```
## [1] 10677
```

```
# Q4  
# How many different users are in the edx dataset?
```

```
uniqueUsers <- n_distinct(edx$userId)  
print(uniqueUsers)
```

```
## [1] 69878
```

```
# Q5  
# How many movie ratings are in each of the  
# following genres in the edx dataset?  
List_of_genre <- c('Drama', 'Comedy', 'Thriller', 'Romance')  
Nbr_of_genres <- sapply(List_of_genre, function(g){  
  edx %>% filter(str_detect(genres, g)) %>% tally()  
})
```

```
edx %>% separate_rows(genres, sep = "\\|") %>%
  group_by(genres) %>%
  summarize(count = n())
```

```
## # A tibble: 20 x 2
##   genres      count
##   <chr>      <int>
## 1 (no genres listed)      7
## 2 Action      2560545
## 3 Adventure   1908892
## 4 Animation   467168
## 5 Children    737994
## 6 Comedy      3540930
## 7 Crime       1327715
## 8 Documentary   93066
## 9 Drama       3910127
## 10 Fantasy     925637
## 11 Film-Noir   118541
## 12 Horror      691485
## 13 IMAX        8181
## 14 Musical     433080
## 15 Mystery     568332
## 16 Romance    1712100
## 17 Sci-Fi      1341183
## 18 Thriller    2325899
## 19 War         511147
## 20 Western     189394
```

```
# Q6
# Which movie has the greatest number of ratings?
N_Ratings <- edx %>% group_by(movieId) %>%
  summarize(N_Ratings = n(), movieTitle = first(title)) %>%
  arrange(desc(N_Ratings)) %>%
  top_n(10, N_Ratings)

print(N_Ratings)
```

```
## # A tibble: 10 x 3
##   movieId N_Ratings movieTitle
##   <int>   <int> <chr>
## 1     296    31362 Pulp Fiction (1994)
## 2     356    31079 Forrest Gump (1994)
## 3     593    30382 Silence of the Lambs, The (1991)
## 4     480    29360 Jurassic Park (1993)
## 5     318    28015 Shawshank Redemption, The (1994)
## 6     110    26212 Braveheart (1995)
## 7     457    25998 Fugitive, The (1993)
## 8     589    25984 Terminator 2: Judgment Day (1991)
## 9     260    25672 Star Wars: Episode IV - A New Hope (a.k.a. Star Wars) (197~
## 10    150    24284 Apollo 13 (1995)
```

```
# Q7
# What are the five most given ratings in order from most to least?
N_Ratings <- edx %>% group_by(rating) %>%
  summarize(number = n())

N_Ratings %>% top_n(5) %>% arrange(desc(number))
```

```
## Selecting by number
```

```
## # A tibble: 5 x 2
##   rating number
##   <dbl>   <int>
## 1     4 2588430
## 2     3 2121240
## 3     5 1390114
## 4   3.5  791624
## 5     2  711422
```

```
# Q8
# True or False: In general, half star ratings are less common than whole star
# ratings (e.g., there are fewer ratings of 3.5 than there are ratings of 3 or 4, etc.).
```

```
N_Ratings %>%
  mutate(halfStar = rating %% 1 == 0.5) %>%
  group_by(halfStar) %>%
  summarize(number = sum(number))
```

```
## # A tibble: 2 x 2
##   halfStar number
##   <lgl>   <int>
## 1 FALSE   7156885
## 2 TRUE    1843170
```

```
#####End Quiz#####
```

2.2 Data Exploration

Exploring the Data Set (EDA)

```
### Before we go any further (build the model), it is important
### to understand the schema of the data, distribution of ratings
### and the relationship of the predictors. Hence my journey toward a better model.
```

```
str(edx)
```

```
## 'data.frame': 9000055 obs. of 6 variables:
## $ userId : int 1 1 1 1 1 1 1 1 1 1 ...
## $ movieId : int 122 185 292 316 329 355 356 362 364 370 ...
## $ rating : num 5 5 5 5 5 5 5 5 5 5 ...
## $ timestamp: int 838985046 838983525 838983421 838983392 838983392 838984474 838983653 838984885 8...
## $ title : chr "Boomerang (1992)" "Net, The (1995)" "Outbreak (1995)" "Stargate (1994)" ...
## $ genres : chr "Comedy|Romance" "Action|Crime|Thriller" "Action|Drama|Sci-Fi|Thriller" "Action|A...
```

```
# From the initial of the EDA, we notice that edx has 6 variables described as follow:
```

```
# • "user_id : int" : a unique identifier of the user who made the rating
```

```
# • "movie_id : int " : a unique identifier of the rated movie
```

```
# • "rating : num " : the score of the rating on a five-star scale
```

```
# • "timestamp : int": the timestamp of the ratings, represented in seconds since midnight Coordinated
```

```
# • "title : chr": the title of the rated movie with the release year in parentheses
```

```
# • "genres : chr": a sequence of genres to which the rated movie belongs
```

```
# • "year Rated: num": 1996 1996 1996 1996 1996 ...
```

More EDA: Exploratory Data Analysis

Dataset Dimenions

```
###Check Dimensions (rows and columns ) of both final_holdout_test and train tables
cat("\nEdx (it contents):",dim(edx))
```

```
##
## Edx (it contents): 9000055 6
```

```
cat("\nfinal holdout test (it contents):",dim(final_holdout_test))
```

```
##
## final holdout test (it contents): 999999 6
```

```
# The following table shows the schema and content of edx dataset
head(edx)
```

```
##   userId movieId rating timestamp                title
## 1      1     122      5 838985046      Boomerang (1992)
## 2      1     185      5 838983525      Net, The (1995)
## 4      1     292      5 838983421      Outbreak (1995)
## 5      1     316      5 838983392      Stargate (1994)
## 6      1     329      5 838983392 Star Trek: Generations (1994)
## 7      1     355      5 838984474      Flintstones, The (1994)
##                                genres
## 1                      Comedy|Romance
## 2          Action|Crime|Thriller
## 4 Action|Drama|Sci-Fi|Thriller
## 5          Action|Adventure|Sci-Fi
## 6 Action|Adventure|Drama|Sci-Fi
## 7          Children|Comedy|Fantasy
```

```
# Dissecting genres
# 2.2.1 No specific Genres(mix genres) and Genres
# The data set contains 797 different combinations of genres.
# The following table contents the first list of genres.
```

```
edx %>% group_by(genres) %>%
  summarise(n=n()) %>%
  head()
```

```
## # A tibble: 6 x 2
##   genres                                n
##   <chr>                             <int>
## 1 (no genres listed)                   7
## 2 Action                           24482
## 3 Action|Adventure                   68688
## 4 Action|Adventure|Animation|Children|Comedy    7467
## 5 Action|Adventure|Animation|Children|Comedy|Fantasy  187
## 6 Action|Adventure|Animation|Children|Comedy|IMAX    66
```

```
# Here is the second list of genre in an orderly fashions.
tibble(count = str_count(edx$genres, fixed("|")), genres = edx$genres) %>%
  group_by(count, genres) %>%
  summarise(n = n()) %>%
  arrange(-count) %>%
  head()
```

'summarise()' has grouped output by 'count'. You can override using the
'.groups' argument.

```
## # A tibble: 6 x 3
## # Groups:   count [3]
##   count genres                                     n
##   <int> <chr>                                     <int>
## 1     7 Action|Adventure|Comedy|Drama|Fantasy|Horror|Sci-Fi|Thriller 256
## 2     6 Adventure|Animation|Children|Comedy|Crime|Fantasy|Mystery 10975
## 3     6 Adventure|Animation|Children|Comedy|Drama|Fantasy|Mystery 355
## 4     6 Adventure|Animation|Children|Comedy|Fantasy|Musical|Romance 515
## 5     5 Action|Adventure|Animation|Children|Comedy|Fantasy 187
## 6     5 Action|Adventure|Animation|Children|Comedy|IMAX 66
```

```
#Sui generis genre (unique genre)
unique_genre <- edx %>% separate_rows(genres, sep = "\\|") %>%
  group_by(genres) %>%
  summarize(count = n()) %>%
  arrange(desc(count))
print(unique_genre)
```

```
## # A tibble: 20 x 2
##   genres          count
##   <chr>          <int>
## 1 Drama          3910127
## 2 Comedy         3540930
## 3 Action         2560545
## 4 Thriller       2325899
## 5 Adventure      1908892
## 6 Romance        1712100
## 7 Sci-Fi         1341183
## 8 Crime          1327715
## 9 Fantasy         925637
## 10 Children       737994
## 11 Horror         691485
## 12 Mystery        568332
## 13 War            511147
## 14 Animation      467168
## 15 Musical        433080
## 16 Western        189394
## 17 Film-Noir     118541
## 18 Documentary    93066
## 19 IMAX           8181
## 20 (no genres listed) 7
```

```
# Noticing that the dataset displays 20 different genres &
# 7 other movies that have not listed as genres whatsoever
# (7 movies without genres)
```

```
# 2.2.2 Date conversion and rating period of time
```

```
#Convert Timestamp to year
```

```
edx <- mutate(edx, year Rated = year(as_datetime(timestamp)))
head(edx)
```

```
##   userId movieId rating timestamp title
## 1      1      122      5 838985046 Boomerang (1992)
## 2      1      185      5 838983525 Net, The (1995)
## 4      1      292      5 838983421 Outbreak (1995)
## 5      1      316      5 838983392 Stargate (1994)
## 6      1      329      5 838983392 Star Trek: Generations (1994)
## 7      1      355      5 838984474 Flintstones, The (1994)
##                                     genres year Rated
## 1                                     Comedy|Romance 1996
## 2                                     Action|Crime|Thriller 1996
## 4 Action|Drama|Sci-Fi|Thriller 1996
## 5                                     Action|Adventure|Sci-Fi 1996
## 6 Action|Adventure|Drama|Sci-Fi 1996
## 7 Children|Comedy|Fantasy 1996
```



```
# The following code will create an Histogram showing the Rating Distribution Per Year  
# Period of collecting rating that started over the years
```

```
if(!require(ggthemes))  
  install.packages("ggthemes", repos = "http://cran.us.r-project.org")
```

```
## Loading required package: ggthemes
```

```
library(ggthemes)
```

```
if(!require(scales))  
  install.packages("scales", repos = "http://cran.us.r-project.org")
```

```
## Loading required package: scales
```

```
##
```

```
## Attaching package: 'scales'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##      discard
```

```
## The following object is masked from 'package:readr':
```

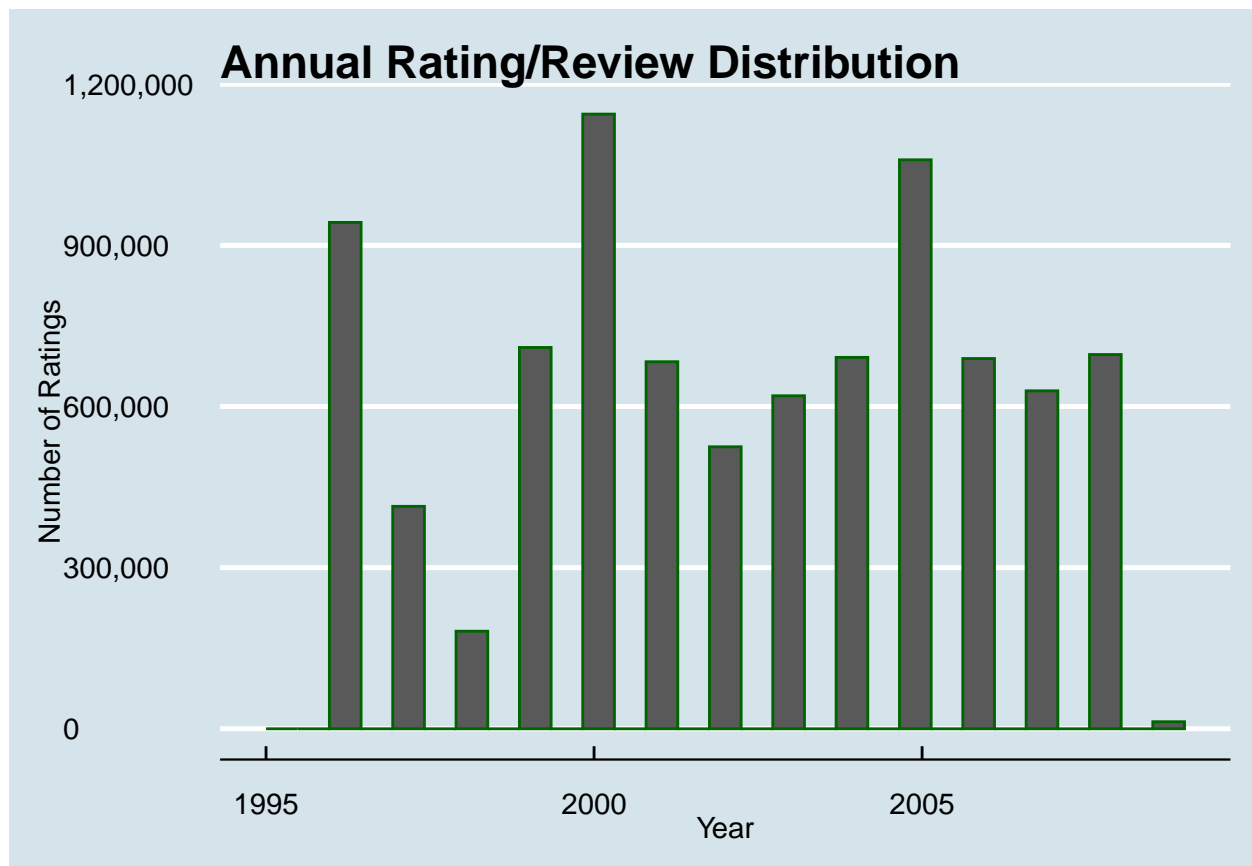
```
##
```

```
##      col_factor
```

```
library(scales)
```

```
edx %>% mutate(year = year(as_datetime(timestamp, origin="1970-01-01"))) %>%  
  ggplot(aes(x=year)) +  
  geom_histogram(color = "dark green") +  
  ggtitle("Annual Rating/Review Distribution") +  
  xlab("Year") +  
  ylab("Number of Ratings") +  
  scale_y_continuous(labels = comma) +  
  theme_economist()
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




#The above plot shows the Rating Distribution Per Year, period which collect rating that started over t

```
#Edx: In depth dissecting (period during which more rating took place)
edx %>% mutate(date = date(as_datetime(timestamp, origin="1970-01-01"))) %>%
  group_by(date, title) %>%
  summarise(count = n()) %>%
  arrange(-count) %>%
  head(15)
```

```
## 'summarise()' has grouped output by 'date'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 15 x 3
## # Groups:   date [4]
##   date      title                                count
##   <date>    <chr>                                <int>
## 1 1998-05-22 Chasing Amy (1997)                    322
## 2 2000-11-20 American Beauty (1999)                277
## 3 1999-12-11 Star Wars: Episode IV - A New Hope (a.k.a. Star Wars) (1977) 254
## 4 1999-12-11 Star Wars: Episode V - The Empire Strikes Back (1980)         251
## 5 1999-12-11 Star Wars: Episode VI - Return of the Jedi (1983)             241
## 6 2005-03-22 Lord of the Rings: The Two Towers, The (2002)                 239
## 7 2005-03-22 Lord of the Rings: The Fellowship of the Ring, The (2001)     227
## 8 2000-11-20 Terminator 2: Judgment Day (1991)                221
## 9 1999-12-11 Matrix, The (1999)                  210
## 10 2000-11-20 Jurassic Park (1993)                201
## 11 2000-11-20 Braveheart (1995)                   197
## 12 2000-11-20 Star Wars: Episode VI - Return of the Jedi (1983)           194
## 13 2000-11-20 Star Wars: Episode IV - A New Hope (a.k.a. Star Wars) (1977) 193
## 14 2000-11-20 Star Wars: Episode V - The Empire Strikes Back (1980)        192
## 15 2005-03-22 Shrek (2001)                      192
```

```
#Graphically speaking let's see which movies have more ratings      than the average mu
mu <- mean(edx$rating)
edx %>% group_by(title) %>%
  summarize(b_i = mean(rating - mu), n = n()) %>% filter(b_i > 0.5, n > 10000) %>%
  ggplot(aes(reorder(title, b_i), b_i, fill = n)) +
  geom_bar(stat = "identity") + coord_flip() + scale_fill_distiller(palette = "PuBuGn") +
  ggtitle("") + xlab("Movie Title") +
  ggtitle("Movie rating - mu,\nfor highest-rated movies that at least 10000 ratings \nor more, wow!!!")
theme_classic()
```



The above graph/plot shows the highest rated movies

Training and Testing Sets:

```
#Training and Testing Sets:
set.seed(2023, sample.kind = "Rounding")

## Warning in set.seed(2023, sample.kind = "Rounding"): non-uniform 'Rounding'
## sampler used

test_index <-createDataPartition(y = edx$rating, times = 1, p = 0.1, list = F)

train_data <-edx[-test_index,]
edx_temp <-edx[test_index,]

#Case where userId and movieId are in both the training and testing sets
test_data <-edx_temp %>% semi_join(train_data, by = "movieId") %>%
  semi_join(train_data, by = "userId")

#Adding removed Rows from the edx_test back into train_data
add_rows_back <-anti_join(edx_temp, test_data)

## Joining with 'by = join_by(userId, movieId, rating, timestamp, title, genres,
## year_rated)'

train_data <-rbind(train_data, add_rows_back)
rm(train_data, test_index, add_rows_back)
```

Cleaning and Analyzing the Data

##Data sui generis (unique). Ensure that data like userIds, movieIds, and genres are not duplicated

```
edx %>% as_tibble()
```

```
## # A tibble: 9,000,055 x 7
##   userId movieId rating timestamp title genres year Rated
##   <int>   <int>   <dbl>      <int> <chr>   <chr>   <dbl>
## 1     1     122     5 838985046 Boomerang (1992) Comed~ 1996
## 2     1     185     5 838983525 Net, The (1995) Actio~ 1996
## 3     1     292     5 838983421 Outbreak (1995) Actio~ 1996
## 4     1     316     5 838983392 Stargate (1994) Actio~ 1996
## 5     1     329     5 838983392 Star Trek: Generations (19~ Actio~ 1996
## 6     1     355     5 838984474 Flintstones, The (1994) Child~ 1996
## 7     1     356     5 838983653 Forrest Gump (1994) Comed~ 1996
## 8     1     362     5 838984885 Jungle Book, The (1994) Adven~ 1996
## 9     1     364     5 838983707 Lion King, The (1994) Adven~ 1996
## 10    1     370     5 838984596 Naked Gun 33 1/3: The Fina~ Actio~ 1996
## # i 9,000,045 more rows
```

```
#Ensure that data are not duplicated (userIds, movieIds, and enres are: sui generis (unique))
edx %>% summarize(unique_users = length(unique(userId)),
                  unique_movies = length(unique(movieId)),
                  unique_genres = length(unique(genres)))
```

```
##   unique_users unique_movies unique_genres
## 1         69878         10677          797
```

```
#####More Ratings: #####
```

```
#Here we go again for more ratings
#Extracting the first date and calculate the age of the movie.
# Find out if the age of the movie effects predicted ratings.
```

```
#Bring the first date to light
```

```
first_date <- stringi::stri_extract(edx$title, regex = "(\\d{4})", comments = TRUE ) %>% as.numeric()
```

```
#Adding the first date
```

```
title_dates <- edx %>% mutate(first_date = first_date)
head(title_dates)
```

```
##   userId movieId rating timestamp title
## 1     1     122     5 838985046 Boomerang (1992)
## 2     1     185     5 838983525 Net, The (1995)
## 4     1     292     5 838983421 Outbreak (1995)
## 5     1     316     5 838983392 Stargate (1994)
## 6     1     329     5 838983392 Star Trek: Generations (1994)
## 7     1     355     5 838984474 Flintstones, The (1994)
##           genres year Rated first_date
## 1 Comedy|Romance 1996 1992
```

```
## 2      Action|Crime|Thriller      1996      1995
## 4 Action|Drama|Sci-Fi|Thriller      1996      1995
## 5      Action|Adventure|Sci-Fi      1996      1994
## 6 Action|Adventure|Drama|Sci-Fi      1996      1994
## 7      Children|Comedy|Fantasy      1996      1994
```

```
#Get rid of timestamp
title_dates <- title_dates %>% select(-timestamp)

head(title_dates)
```

```
##   userId movieId rating      title
## 1      1     122      5 Boomerang (1992)
## 2      1     185      5 Net, The (1995)
## 4      1     292      5 Outbreak (1995)
## 5      1     316      5 Stargate (1994)
## 6      1     329      5 Star Trek: Generations (1994)
## 7      1     355      5 Flintstones, The (1994)
##           genres year Rated first_date
## 1           Comedy|Romance      1996      1992
## 2           Action|Crime|Thriller      1996      1995
## 4 Action|Drama|Sci-Fi|Thriller      1996      1995
## 5           Action|Adventure|Sci-Fi      1996      1994
## 6 Action|Adventure|Drama|Sci-Fi      1996      1994
## 7           Children|Comedy|Fantasy      1996      1994
```

```
#What is the overall mean rating? Here it is:
overall_mean <- mean(edx$rating)
print(overall_mean)
```

```
## [1] 3.512465
```

```
#Now let see if dates are correct
title_dates %>% filter(first_date > 2018) %>% group_by(movieId, title, first_date) %>% summarize(n = n())
```

```
## 'summarise()' has grouped output by 'movieId', 'title'. You can override using
## the '.groups' argument.
```

```
## # A tibble: 6 x 4
## # Groups:   movieId, title [6]
##   movieId title      first_date      n
##   <int> <chr>      <dbl> <int>
## 1     671 Mystery Science Theater 3000: The Movie (1996)      3000    3280
## 2    2308 Detroit 9000 (1973)      9000     22
## 3    4159 3000 Miles to Graceland (2001)      3000    714
## 4    5310 Transylvania 6-5000 (1985)      5000    195
## 5    8864 Mr. 3000 (2004)      3000    146
## 6   27266 2046 (2004)      2046    426
```

```
title_dates %>% filter(first_date < 1900) %>% group_by(movieId, title, first_date) %>% summarize(n = n())
```

```
## 'summarise()' has grouped output by 'movieId', 'title'. You can override using
## the '.groups' argument.
```

```
## # A tibble: 8 x 4
## # Groups:   movieId, title [8]
##   movieId title                                first_date    n
##   <int> <chr>                                <dbl> <int>
## 1   1422 Murder at 1600 (1997)                1600  1566
## 2   4311 Bloody Angels (1732 Høtten: Marerittet Har et Postnu~  1732    9
## 3   5472 1776 (1972)                          1776  185
## 4   6290 House of 1000 Corpses (2003)          1000  367
## 5   6645 THX 1138 (1971)                      1138  464
## 6   8198 1000 Eyes of Dr. Mabuse, The (Tausend Augen des Dr. ~  1000   24
## 7   8905 1492: Conquest of Paradise (1992)     1492  134
## 8  53953 1408 (2007)                          1408  466
```

```
#Now let find the age of the movies
```

```
title_dates <- title_dates %>% mutate(age_of_movie = 2018 - first_date,
                                     rating_date_range = year Rated - first_date)
head(title_dates)
```

```
##   userId movieId rating                                title
## 1     1     122      5                        Boomerang (1992)
## 2     1     185      5                        Net, The (1995)
## 4     1     292      5                        Outbreak (1995)
## 5     1     316      5                        Stargate (1994)
## 6     1     329      5 Star Trek: Generations (1994)
## 7     1     355      5      Flintstones, The (1994)
##                                     genres year Rated first_date age_of_movie
## 1                        Comedy|Romance      1996      1992      26
## 2                Action|Crime|Thriller      1996      1995      23
## 4  Action|Drama|Sci-Fi|Thriller      1996      1995      23
## 5                Action|Adventure|Sci-Fi      1996      1994      24
## 6  Action|Adventure|Drama|Sci-Fi      1996      1994      24
## 7      Children|Comedy|Fantasy      1996      1994      24
##   rating_date_range
## 1                   4
## 2                   1
## 4                   1
## 5                   2
## 6                   2
## 7                   2
```

```
# Skip the graph here.....
```

```
#Rating average: movies, users, average rating by age of movie,      average rating by year
```

```
#Movie rating averages
```

```
movie_avgs <- title_dates %>% group_by(movieId) %>% summarize(avg_movie_rating = mean(rating))
user_avgs <- title_dates %>% group_by(userId) %>% summarize(avg_user_rating = mean(rating))
year_avgs <- title_dates %>% group_by(year Rated) %>% summarize(avg_rating_by_year = mean(rating)) #year
age_avgs <- title_dates %>% group_by(age_of_movie) %>% summarize(avg_rating_by_age = mean(rating)) #age
head(age_avgs)
```



```
## # A tibble: 6 x 2
##   age_of_movie avg_rating_by_age
##   <dbl>         <dbl>
## 1      -6982         2.80
## 2      -2982         2.30
## 3       -982         3.48
## 4        -28         3.73
## 5         8         3.37
## 6        10         3.46
```

```
head(user_avgs)
```

```
## # A tibble: 6 x 2
##   userId avg_user_rating
##   <int>         <dbl>
## 1     1         5
## 2     2        3.29
## 3     3        3.94
## 4     4        4.06
## 5     5        3.92
## 6     6        3.95
```

```
# EDA (exploratory Data Analysis)
#cat("\nTrain set dimension :",dim(edx))
#cat("\nNumber of unique movies :",edx$movieId %>% unique() %>% length())
#cat("\nNumber of unique users :",edx$userId %>% unique() %>% length())
```

```
#Different movies for different genres
cat("\nDifferent movies for different genres :\n")
```

```
##
## Different movies for different genres :
```

```
genres <- c("Drama", "Comedy", "Thriller", "Romance")
sapply(genres, function(g) {
  sum(str_detect(edx$genres, g))
})
```

```
##   Drama   Comedy Thriller  Romance
## 3910127 3540930 2325899 1712100
```

```
#Most rated movies
edx %>% group_by(movieId) %>%
  summarise(n_ratings=n(), title=first(title)) %>%
  top_n(5, n_ratings)
```

```
## # A tibble: 5 x 3
##   movieId n_ratings title
##   <int>     <int> <chr>
## 1     296     31362 Pulp Fiction (1994)
```

```
## 2      318      28015 Shawshank Redemption, The (1994)
## 3      356      31079 Forrest Gump (1994)
## 4      480      29360 Jurassic Park (1993)
## 5      593      30382 Silence of the Lambs, The (1991)
```

```
#Most often ratings (10 top ones)
```


```
edx %>% group_by(rating) %>%
  summarise(n_ratings=n()) %>%
  top_n(10, n_ratings) %>%
  arrange(desc(n_ratings))
```

```
## # A tibble: 10 x 2
##   rating n_ratings
##   <dbl>   <int>
## 1     4     2588430
## 2     3     2121240
## 3     5     1390114
## 4    3.5     791624
## 5     2     711422
## 6    4.5     526736
## 7     1     345679
## 8    2.5     333010
## 9    1.5     106426
## 10   0.5      85374
```

```
#Rating Frequency
```

```
edx %>%
  group_by(rating) %>%
  summarize(count = n()) %>%
  ggplot(aes(x = rating, y = count)) +
  geom_line() +
  ggtitle("Number of frequency/occurrence for each rating")
```

Number of frequency/occurrence for each rating



#Noticing that the most common rating is 4, and the least common is 0.

```
# Way of ratings (number of stars)
way <- ifelse((edx$rating == 1 | edx$rating == 2 | edx$rating == 3 |
              edx$rating == 4 | edx$rating == 5) ,
             "Full_Star",
             "Half_Star")

ratings_way <- data.frame(edx$rating, way)

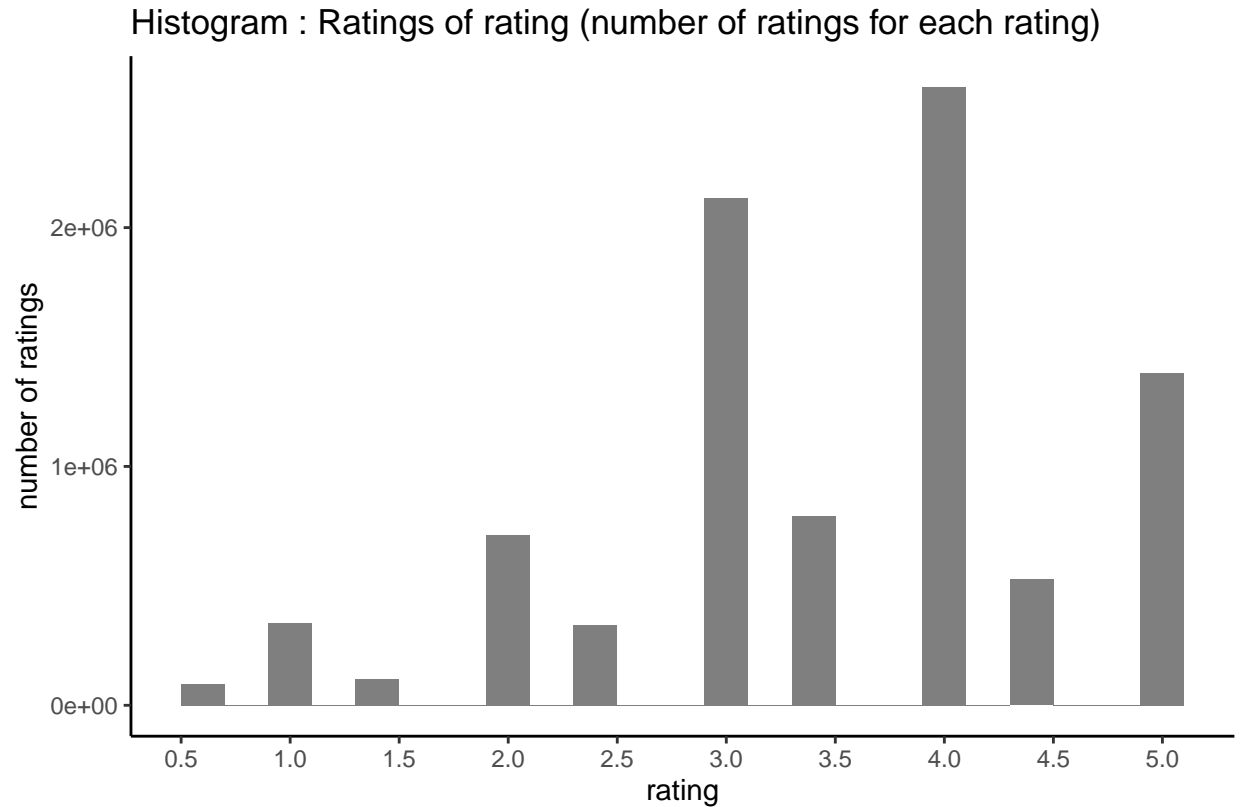
head(ratings_way)
```

```
##   edx.rating    way
## 1         5 Full_Star
## 2         5 Full_Star
## 3         5 Full_Star
## 4         5 Full_Star
## 5         5 Full_Star
## 6         5 Full_Star
```

```
#print(ratings_way)
```

```
#Plotting/Histogram of ratings
ggplot(ratings_way, aes(x= edx.rating, fill = way)) +
  geom_histogram( binwidth = 0.2) +
```


```
scale_x_continuous(breaks=seq(0, 5, by= 0.5)) +
scale_fill_manual(values = c("half_star"="yellow", "full_star"="green")) +
labs(x="rating", y="number of ratings", caption = " According to edx data: set") +
ggtitle("Histogram : Ratings of rating (number of ratings for each rating)")
```



According to edx data: set

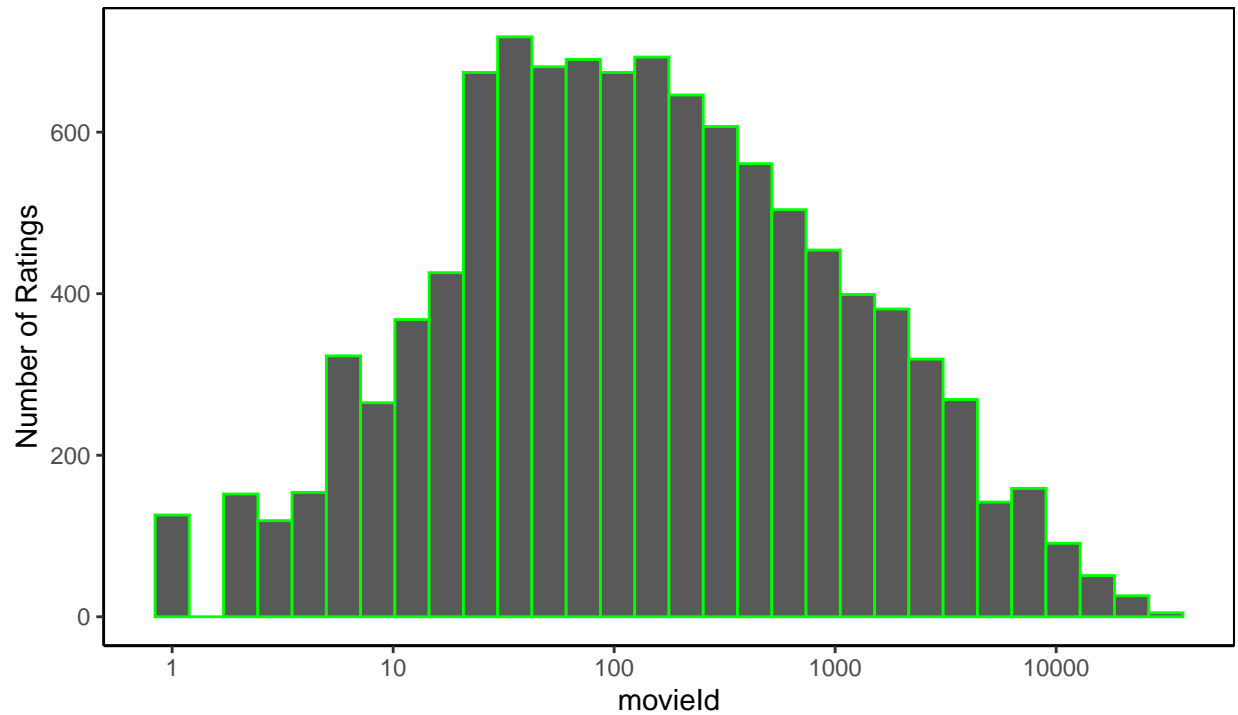
#Plotting/Histogram shows that no zero (0) rating, most ratings are: 4, 3, 5, 3.5 and 2 and the half star ratings are less likely than whole star ratings.

Top Title



#We notice that movies with the highest number of ratings are in the top genres categories such as Jurassic Park(1993), Pulp fiction (1994), Forrest Gump(1994) which are in the top of movie's ratings number, are part of the Drama, Comedy or Action genres. This is what we call blockbusters movies

Visualization: Movies
movies ratings (by movieId)



According to data collection from: edx set

#Noticing that “Reviews” between movies are either less than 1000 or more than 10k. Yes indeed some of them are rated more than others, because many movies are watched by few users and movies like blockbusters have a big impact when it comes to ratings.

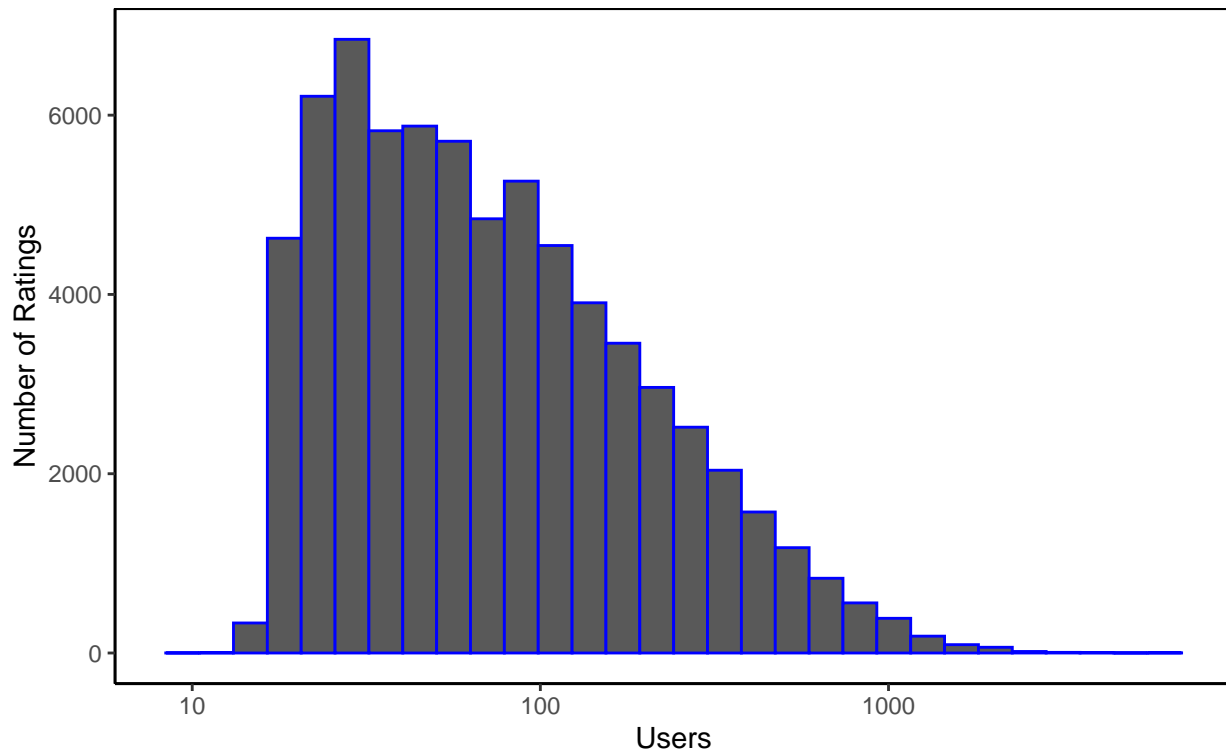
#The following table is example of how most users rate few movies. few users rate more than a thousand

```
edx %>% group_by(userId) %>%  
  summarise(n=n()) %>%  
  arrange(n) %>%  
  head()
```

```
## # A tibble: 6 x 2  
##   userId     n  
##   <int> <int>  
## 1  62516    10  
## 2  22170    12  
## 3  15719    13  
## 4  50608    13  
## 5    901    14  
## 6   1833    14
```

Histogram/Plotting for number of ratings by userId

Visualization: Users
users ratings (by UserId)



Noticing that some users wrote 100 reviews or less, some 1k or more.

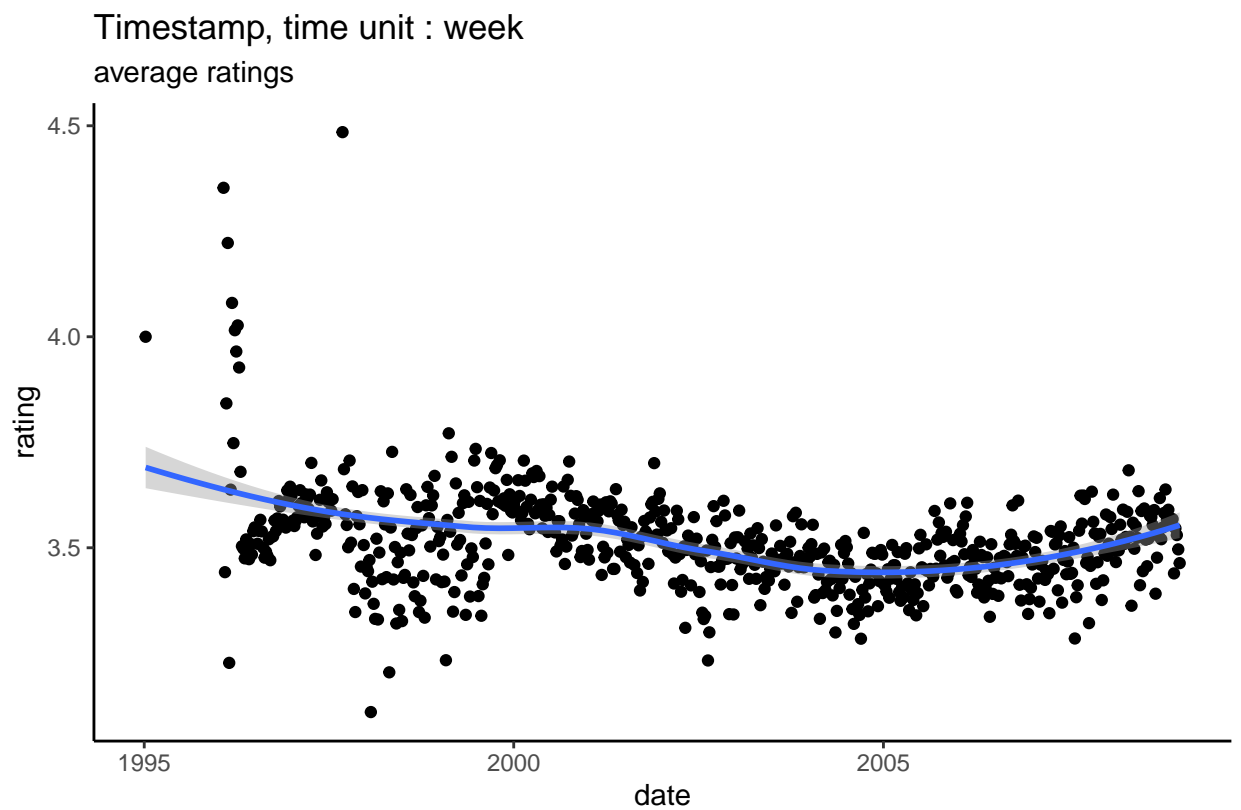
Working with timestamp

#Noticing that the edx set contains the timestamp variable

*#which represents the time and data in which the rating was
#provided. The units (seconds) are important dated back
#January 1, 1970..*

```
edx %>%  
  mutate(date = round_date(as_datetime(timestamp), unit = "week")) %>%  
  group_by(date) %>%  
  summarize(rating = mean(rating)) %>%  
  ggplot(aes(date, rating)) +  
  geom_point() +  
  geom_smooth() +  
  ggtitle("Timestamp, time unit : week")+  
  labs(subtitle = "average ratings",  
       caption = "According to edx: data set")
```

'geom_smooth()' using method = 'loess' and formula = 'y ~ x'



According to edx: data set

Splitting EDX to train and test data

```
### DATA WRANGLING ###

#Splitting EDX to train and test data
set.seed(1998, sample.kind="Rounding")

## Warning in set.seed(1998, sample.kind = "Rounding"): non-uniform 'Rounding'
## sampler used

B <- 100000
edx_sample <- edx[sample(nrow(edx), B, replace = FALSE),]

# Splitting of edx sample on a 80/20 ratio for training purpose
train_index <- createDataPartition(edx_sample$rating, times=1, p=0.8,list=FALSE)
train <- edx_sample[train_index,]
test <- edx_sample[-train_index,]

# edx <- mutate(edx, year_rated = year(as_datetime(timestamp)))

#Wrangling

train$timestamp <- year(as_datetime(train$timestamp))

#extracting release year from title
p <- "(?<=\\(\\d{4}\\)(?=\\))"
train$release_year <- train$title %>% str_extract(p) %>% as.integer()

#Encode the genres column
# str_split(string, p, n = Inf, simplify = FALSE)
# simplify = TRUE
train$genres <- str_split(train$genres, p="\\|")
genre_of_genres <- enframe(train$genres) %>%
  unnest(value) %>%
  mutate(temp = 1) %>%
  pivot_wider(names_from = value, values_from = temp, values_fill = list(temp = 0))
train <- cbind(train, genre_of_genres) %>% select(-name)
train$genres <- NULL

#Adding average rating for each movie by subtracting the total average rating
avg_rating <- mean(train$rating)
movie_score <- train %>% group_by(movieId) %>%
  summarise(movie_score = mean(rating-avg_rating))

#Adding average rating for each user by subtracting the total average rating and movie score
user_score <- train %>% left_join(movie_score, by="movieId") %>%
  mutate(movie_score = ifelse(is.na(movie_score), 0, movie_score)) %>%
  group_by(userId) %>%
  summarise(user_score = mean(rating-avg_rating-movie_score))

train <- train %>% left_join(user_score) %>% left_join(movie_score)
```

```
## Joining with 'by = join_by(userId)'
## Joining with 'by = join_by(movieId)'
```

```
head(train)
```

```
##   userId movieId rating timestamp                                title
## 1   2845    208    3.0    1996                                Waterworld (1995)
## 2  18807   27873    3.5    2005 Metallica: Some Kind of Monster (2004)
## 3   2692   55765    4.0    2008                                American Gangster (2007)
## 4    226   1799    4.0    2003                                Suicide Kings (1997)
## 5  61400   4306    4.0    2008                                Shrek (2001)
## 6  65298   4238    4.0    2008                                Along Came a Spider (2001)
##   year Rated release_year Action Adventure Sci-Fi Documentary Crime Drama
## 1   1996      1995         1         1         1             0         0         0
## 2   2005      2004         0         0         0             1         0         0
## 3   2008      2007         0         0         0             0         1         1
## 4   2003      1997         0         0         0             0         1         1
## 5   2008      2001         0         1         0             0         0         0
## 6   2008      2001         1         0         0             0         1         0
##   Thriller Comedy Mystery Animation Children Fantasy Romance Horror War Musical
## 1         0         0         0             0             0             0         0         0         0
## 2         0         0         0             0             0             0         0         0         0
## 3         1         0         0             0             0             0         0         0         0
## 4         1         1         1             0             0             0         0         0         0
## 5         0         1         0             1             1             1         1         0         0
## 6         1         0         1             0             0             0         0         0         0
##   Western Film-Noir IMAX (no genres listed) user_score movie_score
## 1         0         0         0             0  0.79156046 -0.68357845
## 2         0         0         0             0 -0.38204225 -0.26813727
## 3         0         0         0             0 -0.02195946  0.60686273
## 4         0         0         0             0  0.25960936  0.32801657
## 5         0         0         0             0  0.17085714  0.41043416
## 6         0         0         0             0  0.50443485 -0.08480394
```

```
#Apply same scenario(wrangling) to the test set
```

```
#Convert timestamp to datetime by using only the year
test$timestamp <- year(as_datetime(test$timestamp))
```

```
#extracting release year from title
```

```
p <- "(?<=\\(\\d{4})(?=\\))"
```

```
test$release_year <- test$title %>% str_extract(p) %>% as.integer()
```

```
#Encoding genres column
```

```
test$genres <- str_split(test$genres, p="\\|")
```

```
genre_of_genres <- enframe(test$genres) %>%
```

```
  unnest(value) %>%
```

```
  mutate(temp = 1) %>%
```

```
  pivot_wider(names_from = value, values_from = temp, values_fill = list(temp = 0))
```

```
test <- cbind(test, genre_of_genres) %>% select(-name)
```

```
train$genres <- NULL
```

```

#Adding & removing data (add missing columns of genres that are absent in test set, and
#get rid of those that are not in the train set)
for(col in names(train)){
  if(!col %in% names(test)){
    test$newcol <- 0
    names(test)[names(test)=="newcol"] <- col
  }
}
for(col in names(test)){
  if(!col %in% names(train)){
    test[,col] <- NULL
  }
}

#Average scores on the train set of each movie and user
test$user_score <- NULL
test$movie_score <- NULL
test <- test %>% left_join(user_score, by="userId") %>% left_join(movie_score, by="movieId")

test <- test %>% mutate(user_score = ifelse(is.na(user_score), 0, user_score)) %>% mutate(movie_score =

#Reordering .....
test <- test %>% select(names(train))
head(test)

```

```

##   userId movieId rating timestamp
## 1   4431   1084   4.0      1999
## 2  67235   1210   4.0      2006
## 3  29930   2611   4.5      2005
## 4  37568   7983   4.0      2004
## 5   8739   3094   3.0      2001
## 6   8651   3253   3.0      2004
##
##                                     title year_rated release_year
## 1                                     Bonnie and Clyde (1967)      1999      1967
## 2 Star Wars: Episode VI - Return of the Jedi (1983)      2006      1983
## 3                                     Winslow Boy, The (1999)      2005      1999
## 4                                     Broadway Danny Rose (1984)      2004      1984
## 5                                     Maurice (1987)      2001      1987
## 6                                     Wayne's World (1992)      2004      1992
##   Action Adventure Sci-Fi Documentary Crime Drama Thriller Comedy Mystery
## 1           0           0           0           0           1           1           0           0           0
## 2           1           1           1           0           0           0           0           0           0
## 3           0           0           0           0           0           1           0           0           0
## 4           0           0           0           0           0           0           0           1           0
## 5           0           0           0           0           0           1           0           0           0
## 6           0           0           0           0           0           0           0           1           0
##   Animation Children Fantasy Romance Horror War Musical Western Film-Noir IMAX
## 1           0           0           0           0           0           0           0           0           0           0
## 2           0           0           0           0           0           0           0           0           0           0

```

## 3	0	0	0	0	0	0	0	0	0	0	
## 4	0	0	0	0	0	0	0	0	0	0	
## 5	0	0	0	1	0	0	0	0	0	0	
## 6	0	0	0	0	0	0	0	0	0	0	
##	(no genres listed)			user_score	movie_score						
## 1	0	0.00000000		0.20408495							
## 2	0	0.00000000		0.52464348							
## 3	0	-0.57317073		-0.10904636							
## 4	0	-0.26083520		0.31519606							
## 5	0	-0.05831583		0.28186273							
## 6	0	0.54817874		-0.02722818							

BUILDING of ML MODELS

```
# BUILDING of ML MODELS
```

```
# Baseline comparison .....
```

```
y_hat <- mean(train$rating)
result <- RMSE(test$rating, y_hat)
cat("RMSE :", result)
```

```
## RMSE : 1.051226
```

```
# RMSE is 1.051226 meaning on average the prediction is
# OBOE (off-by-one error) which is not so good
```

```
#Building of ML (Machine Learning) Models
```

```
#Linear Model
```

```
##?
```

```
#Linear Model using the following: timestamp, user_score, release_year, and movie_score
```

```
control <- trainControl(method = "none")
fit_lm <- train(rating~user_score+movie_score+timestamp+
               release_year, data=train, method="lm",
               trControl=control)
print(fit_lm$finalModel)
```

```
##
```

```
## Call:
```

```
## lm(formula = .outcome ~ ., data = dat)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      user_score      movie_score      timestamp      release_year
##    4.7574378        1.0003619        0.9757425       -0.0001211       -0.0005009
```

```
y_hat <- predict(fit_lm, test)
result2 <- RMSE(test$rating, y_hat)
cat("RMSE :", result2)
```

```
## RMSE : 1.042983
```

```
# So let's see how is it for LM without movieId, userId & title
```

```
t2 <- train %>% select(-c("userId", "movieId", "title"))
control <- trainControl(method = "none")
fit_lm <- train(rating~., data=t2, method="lm", trControl=control)
print(fit_lm$finalModel)
```

```
##
```

```
## Call:
```

```
## lm(formula = .outcome ~ ., data = dat)
```

```
##
## Coefficients:
##          (Intercept)                timestamp
##          5.1836910                -0.0002859
##          yearRated                release_year
##          NA                        -0.0005509
##          Action                    Adventure
##          -0.0151763                -0.0023558
##          '\\Sci-Fi\\'              Documentary
##          -0.0009051                0.0479658
##          Crime                     Drama
##          0.0141191                0.0158736
##          Thriller                  Comedy
##          0.0005888                0.0010364
##          Mystery                   Animation
##          0.0048191                0.0219648
##          Children                  Fantasy
##          -0.0339020                0.0136872
##          Romance                   Horror
##          -0.0126802                0.0049235
##          War                       Musical
##          -0.0043971                0.0082115
##          Western                   '\\Film-Noir\\'
##          -0.0054717                0.0134549
##          IMAX '\\(no genres listed)\\'
##          -0.0417882                -0.0287400
##          user_score                movie_score
##          1.0008362                0.9683767
```

```
y_hat <- predict(fit_lm, test)
result3 <- RMSE(test$rating, y_hat)
cat("RMSE :", result3)
```

```
## RMSE : 1.042491
```

```
# Now LM without movieId, userId & title
t2 <- train %>% select(-c("userId", "movieId", "title"))
control <- trainControl(method = "none")
fit_lm <- train(rating~., data=t2, method="lm", trControl=control)
print(fit_lm$finalModel)
```

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Coefficients:
##          (Intercept)                timestamp
##          5.1836910                -0.0002859
##          yearRated                release_year
##          NA                        -0.0005509
##          Action                    Adventure
##          -0.0151763                -0.0023558
##          '\\Sci-Fi\\'              Documentary
```

```
##          -0.0009051          0.0479658
##          Crime          Drama
##          0.0141191          0.0158736
##          Thriller          Comedy
##          0.0005888          0.0010364
##          Mystery          Animation
##          0.0048191          0.0219648
##          Children          Fantasy
##          -0.0339020          0.0136872
##          Romance          Horror
##          -0.0126802          0.0049235
##          War          Musical
##          -0.0043971          0.0082115
##          Western          '\\Film-Noir\\'
##          -0.0054717          0.0134549
##          IMAX '\\(no genres listed)\\'
##          -0.0417882          -0.0287400
##          user_score          movie_score
##          1.0008362          0.9683767
```

```
y_hat <- predict(fit_lm, test)
result3 <- RMSE(test$rating, y_hat)
cat("RMSE :", result3)
```

```
## RMSE : 1.042491
```

#RMSE = 1.042491. It is slightly better than the baseline score, but not good enough.

#Decision Tree

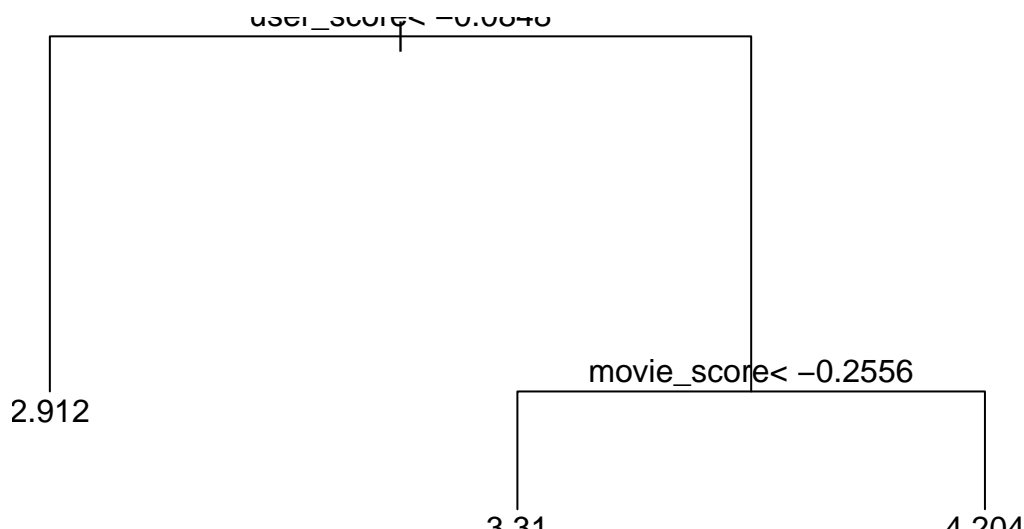
```
fit_tree <- train(rating~user_score+movie_score+timestamp+release_year, data=train, method="rpart")
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

```
print(fit_tree$results)
```

```
##          cp          RMSE Rsquared          MAE          RMSESD RsquaredSD          MAESD
## 1 0.07660190 0.8463595 0.3588611 0.6614992 0.02436197 0.037277092 0.01072485
## 2 0.08038463 0.8864266 0.2967925 0.6762078 0.02470256 0.039419851 0.01075139
## 3 0.24239825 0.9872902 0.2391572 0.7683068 0.07022719 0.002803921 0.08427557
```

```
plot(fit_tree$finalModel)
text(fit_tree$finalModel)
```



*# Noticing that prediction using Decision Tree is way too simple, with
only 3 predicted ratings and few conditions.*

```
y_hat <- predict(fit_tree, test)
result4 <- RMSE(test$rating, y_hat)
cat("RMSE :", result4)
```

```
## RMSE : 1.062504
```

*#The outcome is even worse than the baseline model. So let
#try the linear model with regularization*

#Linear Model with regularizationwith only user_score and movie_score

```
#splitting the train set into 2 to calculate the best of the two
indx <- createDataPartition(train$rating, times=1, p=0.8, list=FALSE)
tpart_1 <- train[indx, ]
tpart_2 <- train[-indx, ]
```

```
#calculating the best ones
best_ones <- seq(0, 10, 0.25)
rmsees <- sapply(best_ones, function(l){
  avg_rating <- mean(tpart_1$rating)
  movie_score <- tpart_1 %>%
```



```

    group_by(movieId) %>%
    summarize(b_m = sum(rating - avg_rating)/(n()+1))
user_score <- tpart_1 %>%
  left_join(movie_score, by="movieId") %>%
  mutate(b_m = ifelse(is.na(b_m), 0, b_m)) %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_m - avg_rating)/(n()+1))
predicted_ratings <-
  tpart_2 %>%
  left_join(movie_score, by = "movieId") %>%
  left_join(user_score, by = "userId") %>%
  mutate(b_m = ifelse(is.na(b_m), 0, b_m)) %>%
  mutate(b_u = ifelse(is.na(b_u), 0, b_u)) %>%
  mutate(pred = avg_rating + b_m + b_u) %>%
  .$pred
return(RMSE(predicted_ratings, tpart_2$rating))
})

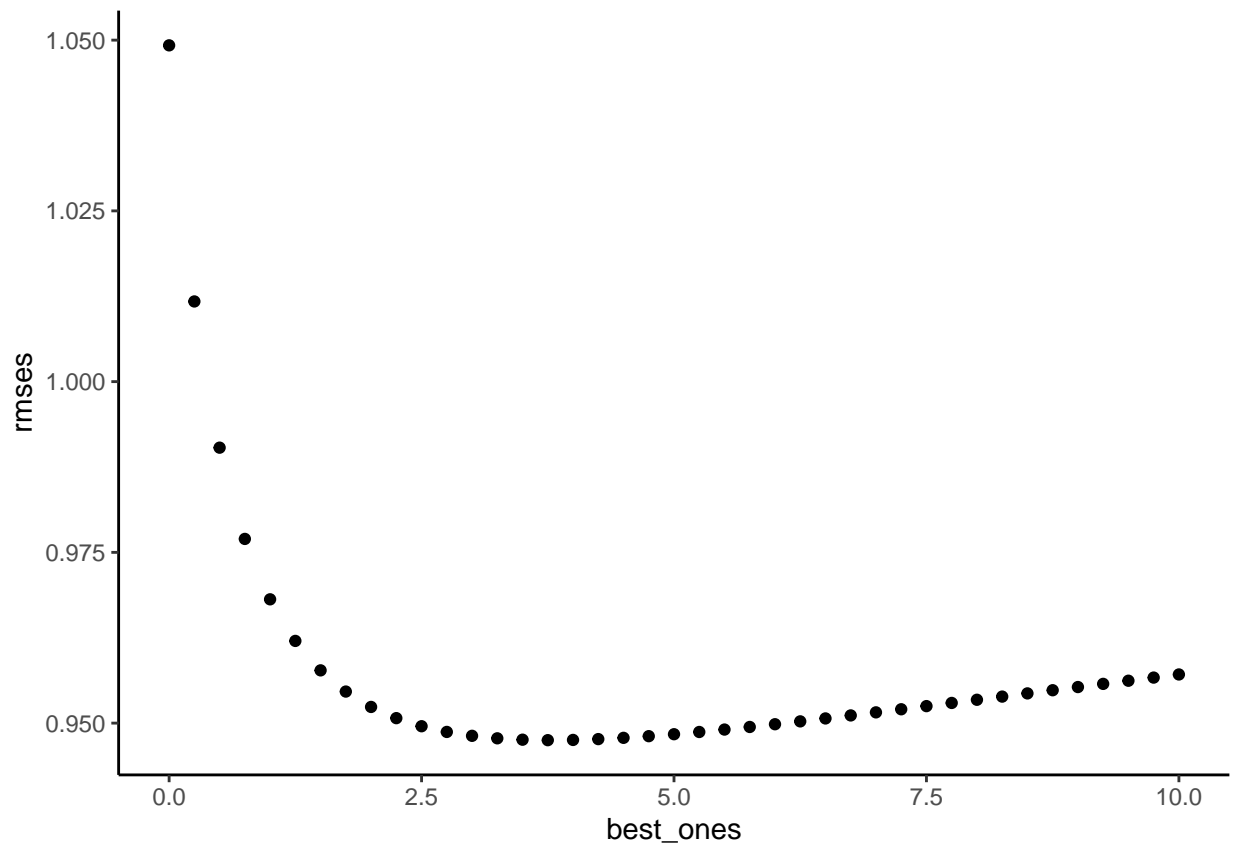
b1 <- best_ones[which.min(rmses)]
qplot(best_ones, rmses)

```

```

## Warning: 'qplot()' was deprecated in ggplot2 3.4.0.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.

```



```
print(b1)
```

```
## [1] 3.75
```

```
# b1 = 3.75
```

```
# The b1 which minimizes the RMSE is 3.75, so let use it to  
#train the model and predict the test set
```

```
#Prediction
```

```
b1 <- 3.75
```

```
avg_rating <- mean(train$rating)
```

```
movie_score <- train %>%
```

```
  group_by(movieId) %>%
```

```
  summarize(b_m = sum(rating - avg_rating)/(n()+b1))
```

```
user_score <- train %>%
```

```
  left_join(movie_score, by="movieId") %>%
```

```
  mutate(b_m = ifelse(is.na(b_m), 0, b_m)) %>%
```

```
  group_by(userId) %>%
```

```
  summarize(b_u = sum(rating - b_m - avg_rating)/(n()+b1))
```

```
predicted_ratings <-
```

```
  test %>%
```

```
  left_join(movie_score, by = "movieId") %>%
```

```
  left_join(user_score, by = "userId") %>%
```

```
  mutate(b_m = ifelse(is.na(b_m), 0, b_m)) %>%
```

```
  mutate(b_u = ifelse(is.na(b_u), 0, b_u)) %>%
```

```
  mutate(pred = avg_rating + b_m + b_u) %>%
```

```
  .$pred
```

```
result5 <- RMSE(test$rating, predicted_ratings)
```

```
cat("RMSE :", result5)
```

```
## RMSE : 0.9383446
```

```
# The result of the RMSE < 1, which is much better than  
# the other models. Now let use the genres columns to see how  
# the predictions will be. So what is the effect of  
# genre on the ratings?
```

```
not_genres <- c("userId", "movieId", "rating", "timestamp", "title", "release_year", "user_score", "mov
```

```
genres <- colnames(train)[!colnames(train) %in% not_genres]
```

```
genres
```

```
## [1] "year Rated"      "Action"          "Adventure"  
## [4] "Sci-Fi"          "Documentary"     "Crime"  
## [7] "Drama"           "Thriller"        "Comedy"  
## [10] "Mystery"         "Animation"       "Children"  
## [13] "Fantasy"         "Romance"         "Horror"  
## [16] "War"             "Musical"         "Western"  
## [19] "Film-Noir"      "IMAX"            "(no genres listed)"
```

```

#What is the average ratings for each genre? Let see.
genre_scores <- data.frame(genre="",m=0, sd=0)
for(genre in genres){
  results <- train %>% filter(train[colnames(train)==genre]==1) %>%
    summarise(m=mean(rating), sd=sd(rating))
  genre_scores <- genre_scores %>% add_row(genre=genre, m=results$m, sd=results$sd)
}

```

```

## Warning: Using one column matrices in 'filter()' was deprecated in dplyr 1.1.0.
## i Please use one dimensional logical vectors instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.

```

```

genre_scores <- genre_scores[-1,]
genre_scores[is.na(genre_scores)] <- 0
genre_scores

```

```

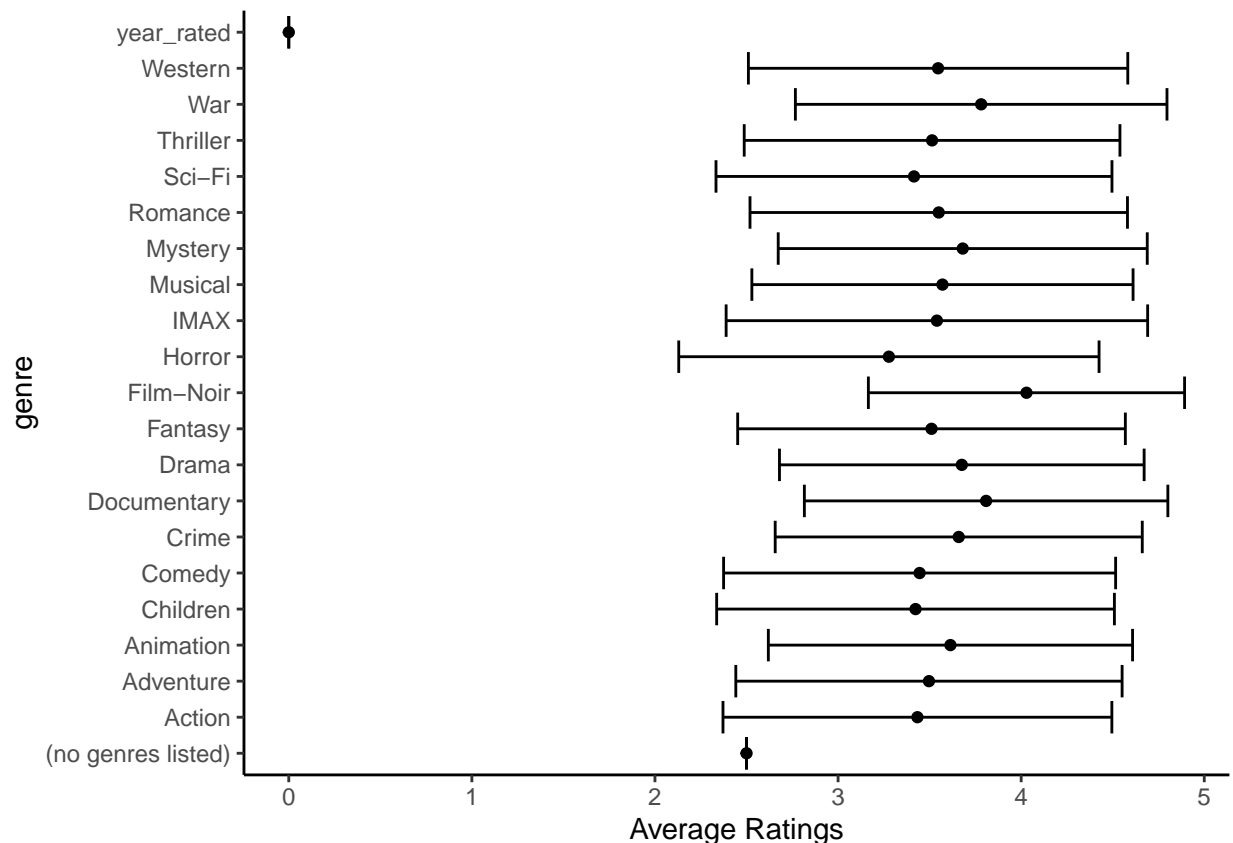
##           genre           m           sd
## 2      year Rated 0.000000 0.0000000
## 3      Action 3.433409 1.0619524
## 4      Adventure 3.496405 1.0548618
## 5      Sci-Fi 3.414611 1.0812908
## 6      Documentary 3.808537 0.9926102
## 7      Crime 3.658984 1.0023646
## 8      Drama 3.675710 0.9958227
## 9      Thriller 3.513454 1.0260652
## 10     Comedy 3.445359 1.0705303
## 11     Mystery 3.680841 1.0078153
## 12     Animation 3.613767 0.9945636
## 13     Children 3.423029 1.0861241
## 14     Fantasy 3.510395 1.0586236
## 15     Romance 3.549792 1.0307855
## 16     Horror 3.277596 1.1477909
## 17     War 3.781634 1.0147717
## 18     Musical 3.570237 1.0409776
## 19     Western 3.546312 1.0360351
## 20     Film-Noir 4.029081 0.8635001
## 21      IMAX 3.539683 1.1510920
## 22 (no genres listed) 2.500000 0.0000000

```

```

#Plotting genres
genre_scores %>% ggplot(aes(x=m, y=genre)) +
  geom_point() +
  xlab("Average Ratings") +
  geom_errorbarh(aes(xmin=m-sd, xmax=m+sd))

```



*# In the plot we notice that different genres have different
ratings average. So it is ideal to use the average of the genres
of a movie to predict the ratings if the movie and the user
in the test are not seen in the training set.*

*#The actual issue with the regularized model is if there
is a case in the test data that has new movie and new user,
the model can only predict with the average of all ratings.
With the added feature of genres, it will probably change the landscape for better
Minimize the RMSE.*

#Regularized model with genre feature

```
b1 <- 3.75
avg_rating <- mean(train$rating)
movie_score <- train %>%
  group_by(movieId) %>%
  summarize(b_m = sum(rating - avg_rating)/(n()+b1))
user_score <- train %>%
  left_join(movie_score, by="movieId") %>%
  mutate(b_m = ifelse(is.na(b_m), 0, b_m)) %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_m - avg_rating)/(n()+b1))
genre_score <- as.matrix(test[, genres]) %*% genre_scores$m
n_genres <- rowSums(test[,genres])
genre_score <- genre_score / n_genres
```

```

#What is the effect of using the genre_scores if the user and
#movie are unknown?
predicted_ratings <-
  test %>%
    left_join(movie_score, by = "movieId") %>%
    left_join(user_score, by = "userId") %>%
    cbind(genre_score) %>%
    mutate(pred = genre_score) %>%
    mutate(pred = ifelse(!is.na(b_m) || !is.na(b_u),
                        avg_rating + replace_na(b_m,0) + replace_na(b_u,0),
                        pred))

result6 <- RMSE(test$rating, predicted_ratings$pred)
cat("RMSE :", result6)

```

```
## RMSE : 0.9491269
```

```

# We notice that the improvement is very slim in RMSE
# when using genres for prediction purposes

```

```
#Put them together (table of RMSE Results)
```

```

data.frame(
  method=c("Naive Prediction", "Linear Model (with 4 features)", "Linear Model (with all features)", "D
  rmse=c(result, result2, result3, result4, result5, result6))

```

```

##
## 1
## 2
## 3
## 4
## 5 Linear Model with Regularisation(only using movie and user scores)
## 6 Linear Model with Regularisation(movie, user, and genre scores)

```

	method	rmse
1	Naive Prediction	1.0512258
2	Linear Model (with 4 features)	1.0429830
3	Linear Model (with all features)	1.0424912
4	Decision Tree	1.0625044
5	Linear Model with Regularisation(only using movie and user scores)	0.9383446
6	Linear Model with Regularisation(movie, user, and genre scores)	0.9491269

Using the model on the final_holdout_test data

```
#Training the final model

# Now let apply the final model to validation set
# After dissecting to come up with different models, we can use
# the best performing model in the previous section, which is the regularized model.
# The validation data is set in a way so the users and movies in the data are all present in the
# edx data.
# So, it is not recommended using the following:
#genre fech has both unknown user and unknown movie, not the best choice

#Using the model on the final_holdout_test data
b1 <- 3.75
avg_rating <- mean(edx$rating)
movie_score <- edx %>%
  group_by(movieId) %>%
  summarize(b_m = sum(rating - avg_rating)/(n()+b1))
user_score <- edx %>%
  left_join(movie_score, by="movieId") %>%
  mutate(b_m = ifelse(is.na(b_m), 0, b_m)) %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_m - avg_rating)/(n()+b1))
predicted_ratings <-
  final_holdout_test %>%
  left_join(movie_score, by = "movieId") %>%
  left_join(user_score, by = "userId") %>%
  mutate(b_m = ifelse(is.na(b_m), 0, b_m)) %>%
  mutate(b_u = ifelse(is.na(b_u), 0, b_u)) %>%
  mutate(pred = avg_rating + b_m + b_u) %>%
  .$pred

final_result <- RMSE(final_holdout_test$rating, predicted_ratings)
cat("RMSE :", final_result)
```

```
## RMSE : 0.8648477
```

The final RMSE is 0.8648477 which is the required RMSE (0.8649) to get the maximum point for the EDX Capstone MovieLens projects.

Conclusion

This MovieLens dataset model provided by edx is a beautiful project to work on. Indeed outcomes can be quite different, close to perfect meaning RMSE could almost be second to none if average raters(users) don't play bias, meaning user doesn't rate a particularly good/popular movie with a large margin bi, and vice versa. Probabilistically speaking if my statement about RMSE is not impossible, then it has to be probable, even if the chance is infinitesimally small.

Polymorphically speaking, I used quite a few machine learning algorithms to come up with predictions movie ratings for this MovieLens dataset. As expected results are different from one another for RMSE. Among those steps (algorithms) used. The regularized model would be ideal with the users side effect to lower RMSE.

Is there a better way or much room for improvement in a nutshell? Off course yes, because life is polymorphic but I haven't come across it yet.

Hence this model can be improved by adding other accoutrement(age, year, genre,...) and on how users should/could rate movies. Also we can apply different machine learning models to improve, hence a better polished outcome.